

From Code to Containers: Understanding Serverless Beyond Lambda

Avinash Dalvi

Head of technology, Nushift Technologies

AWS User group Leader, Bengaluru

AWS Community Builder



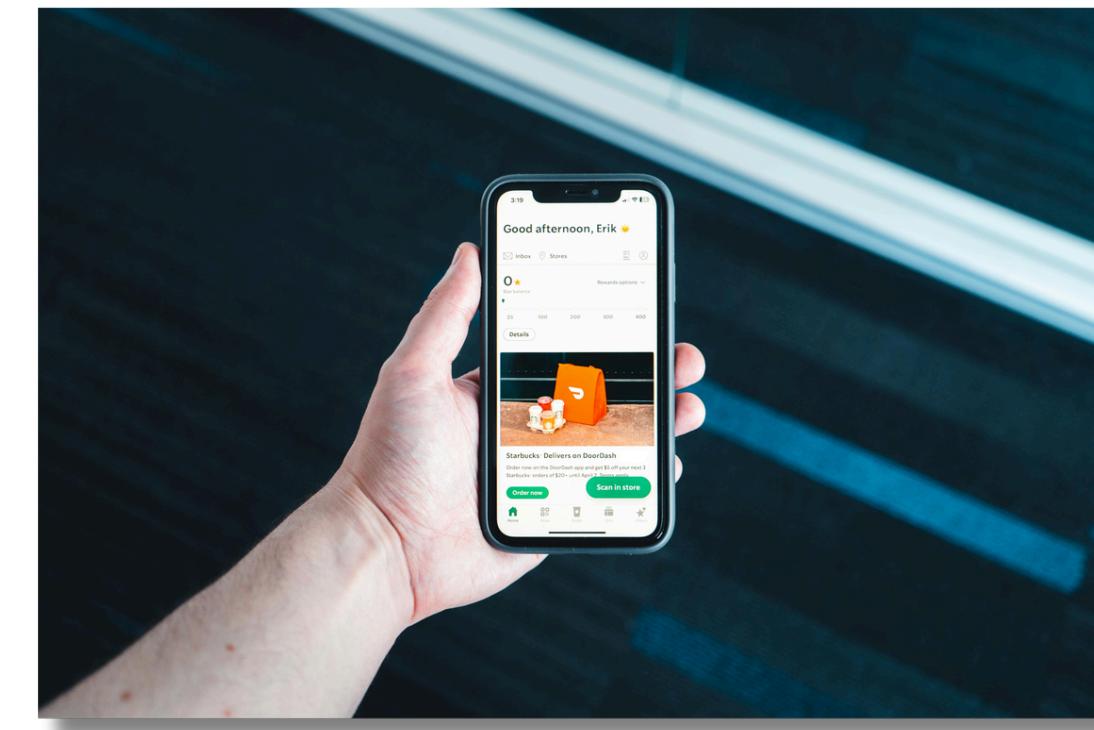
It's Friday night,
you and your
friends want
pizza...



Option 1: Make
it yourself



Option 2: Just
order pizza



Which would you choose? 🤔

Traditional Servers (EC2)

- ✓ You manage the "kitchen" (server)
- ✓ You maintain the "oven" (patching, updates)
- ✓ You pay for it 24/7 (idle or busy)
- ✓ Full control, full responsibility

Serverless

- ✓ Just bring your code (order pizza)
- ✓ Specify requirements (toppings, size, base)
- ✓ AWS handles everything else
- ✓ Pay only for execution time**

How do you "order pizza" in AWS?



Serverless \neq No Servers

Serverless = You don't worry
about servers!

Ordering Your Serverless Pizza



BASE(RUNTIME/LANGUAGE)

- Python, Node.js, Java, Go, .NET, Ruby...
- Or bring your own custom base (containers)



TOPPINGS (RESOURCES)

- RAM: 128MB to 10GB+
- Storage: Temporary disk space
- CPU: Scales with memory



SIZE (CODE COMPLEXITY)

- Small: Simple functions
- Medium: Moderate logic
- Large: Complex applications



WHEN SHOULD IT RUN?

- Immediate (event-triggered)
- Scheduled (cron jobs)
- On-demand (API calls)

The Serverless Menu



Both are "ordering pizza" but different options...

Different options...

Aspect	Lambda	Fargate
What it is	Standard menu	Custom recipe
Base options	Pre-set runtimes	Any container image
Customization	Limited to menu	Fully customizable
Execution time	Max 15 minutes	Unlimited
Code size	250MB unzipped	No limit
Use case	Quick functions	Long processes
Cold starts	Sometimes	More consistent

Basic settings [Info](#)

Description - optional

Memory [Info](#)

Your function is allocated CPU proportional to the memory configured.

128 MB

Set memory to between 128 MB and 10240 MB

Ephemeral storage [Info](#)

You can configure up to 10 GB of ephemeral storage (/tmp) for your function. [View pricing](#)

512 MB

Set ephemeral storage (/tmp) to between 512 MB and 10240 MB.

SnapStart [Info](#)

Reduce startup time by having Lambda cache a snapshot of your function after the function has initialized. To evaluate whether your function code is resilient to snapshot operations, review the [SnapStart compatibility considerations](#). For Python and .NET runtimes, [view pricing](#).

None



Supported runtimes: .NET 8 (C#/F#/PowerShell), Java 11, Java 17, Java 21, Python 3.12, Python 3.13.

Timeout

1 min 0 sec

Execution role

Choose a role that defines the permissions of your function. To create a custom role, go to the [IAM console](#).

Use an existing role

Create a new role from AWS policy templates

Existing role

Choose an existing role that you've created to be used with this Lambda function. The role must have permission to upload logs to Amazon CloudWatch Logs.

LambdaChime



[View the LambdaChime role](#) on the IAM console.

Image Resizer

```
1 # Your code (the pizza recipe)
2 import boto3
3 from PIL import Image
4
5 def lambda_handler(event, context):
6     # Get uploaded image from S3
7     bucket = event['Records'][0]['s3']['bucket']['name']
8     key = event['Records'][0]['s3']['object']['key']
9
10    # Resize image
11    image = Image.open(f'/tmp/{key}')
12    image.thumbnail((200, 200))
13
14    # Save thumbnail
15    image.save(f'/tmp/thumb_{key}')
16    return {'status': 'success'}
```

When Lambda Hits the Wall

-  Upload video → Lambda works great!
-  Generate thumbnail → Lambda perfect!
-  Transcode video → Lambda fails!

Why?



- Video processing takes 45 minutes
- Lambda timeout: 15 minutes
- Need more control over environment
- Larger dependencies required

Fargate - Bring Your Own Recipe

- You bring a container (Docker image)
- AWS runs it without managing servers



Fargate - The Freedom



MORE CONTROL

- Any programming language
- Any runtime version
- Custom OS dependencies
- Specific tools/libraries



MORE CAPACITY

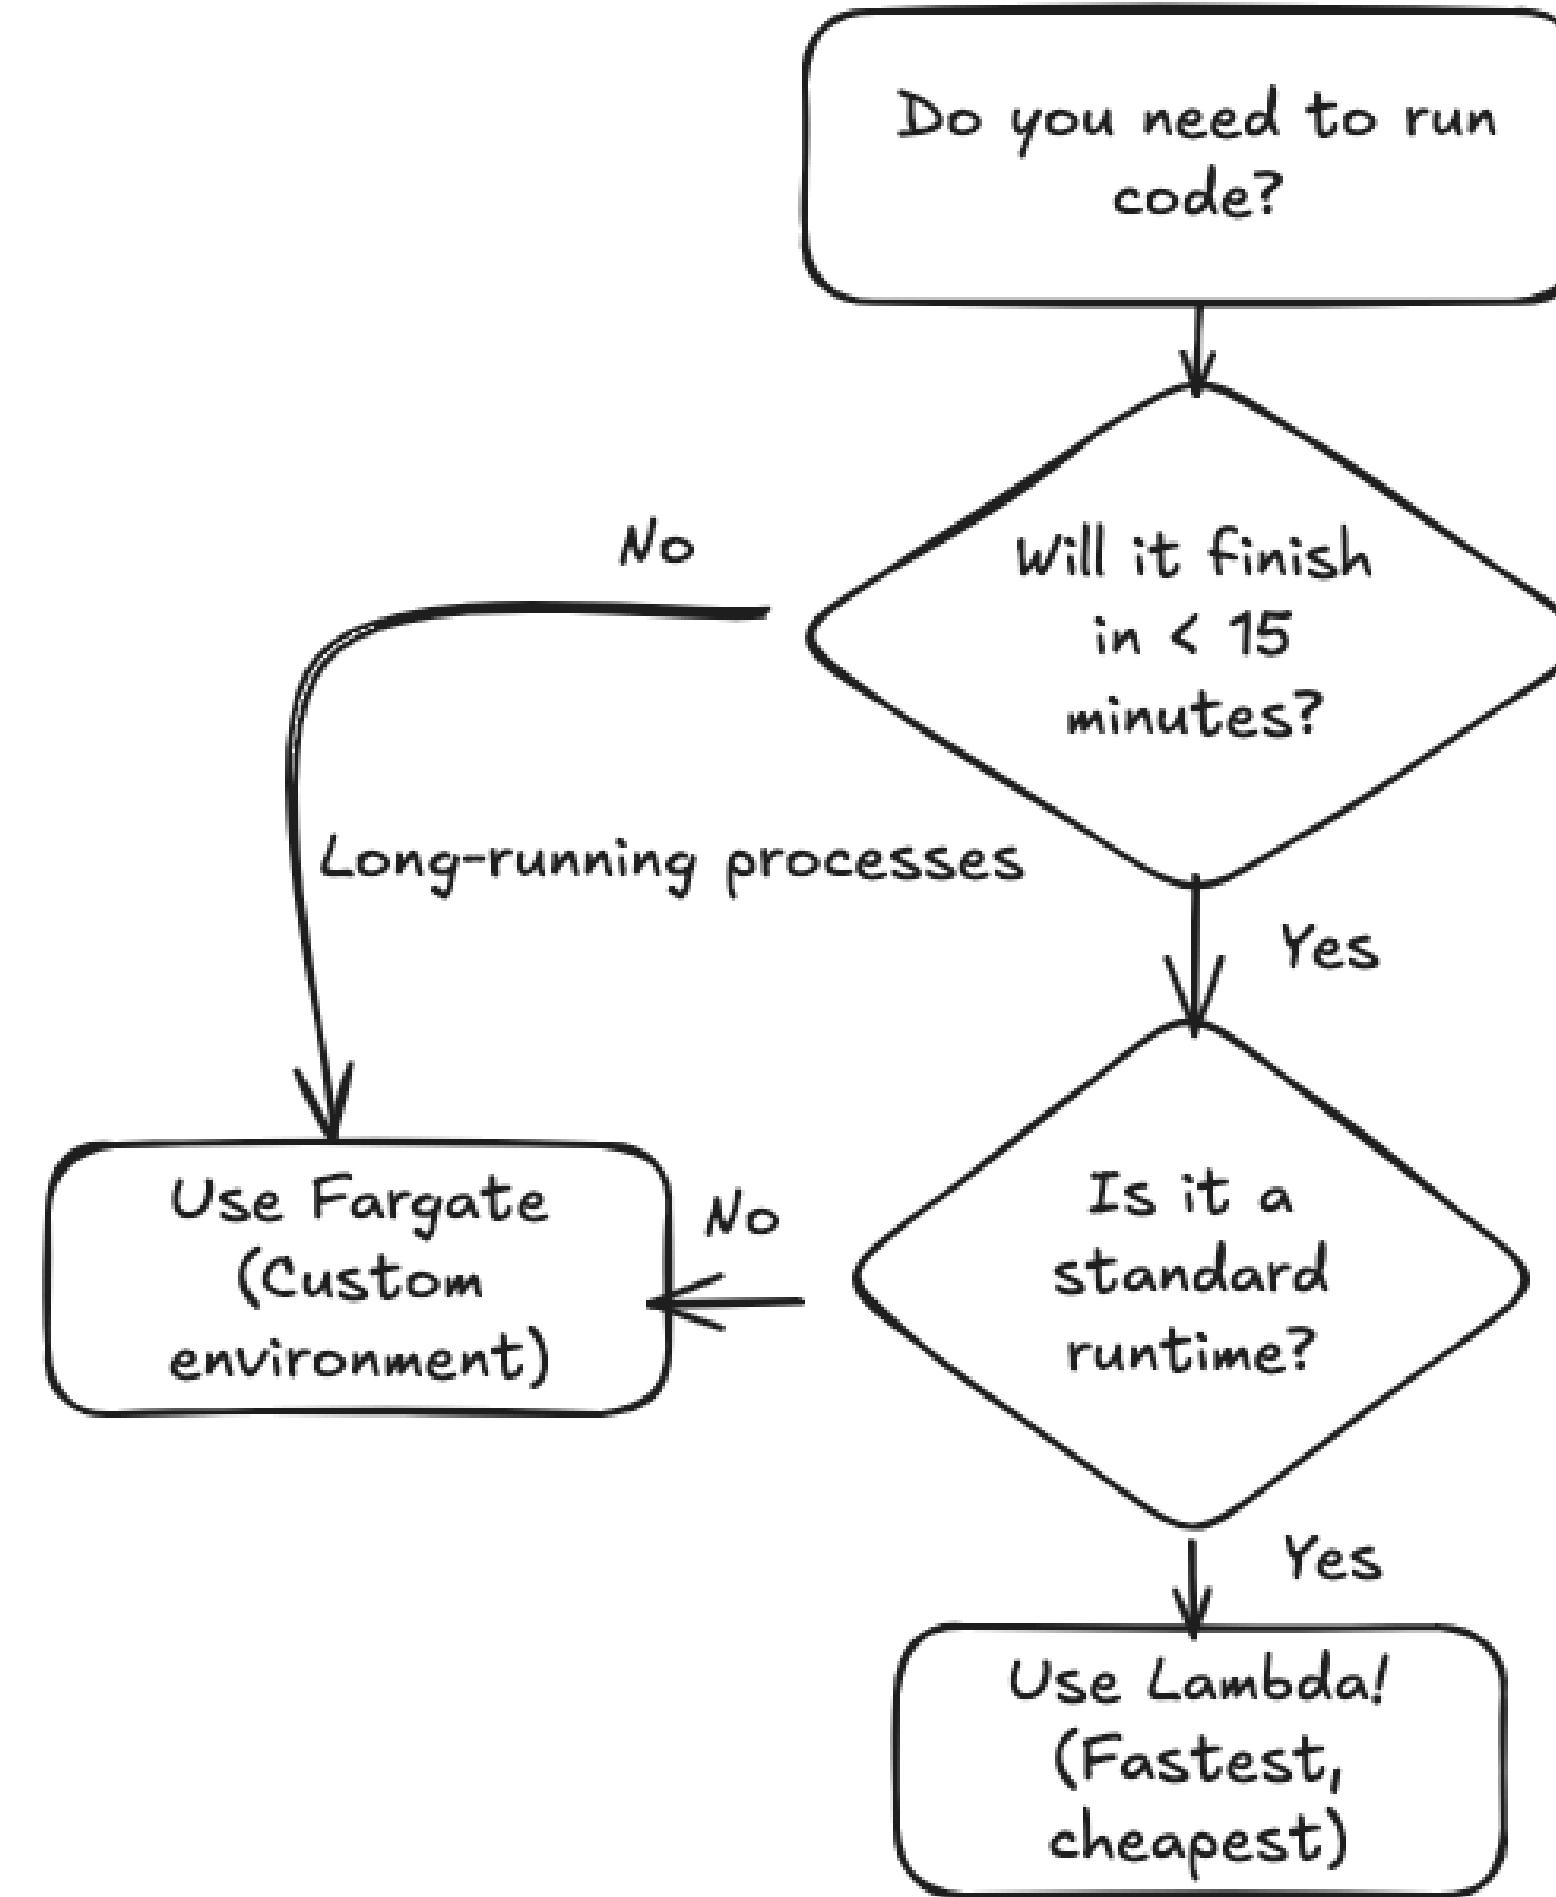
- Run for hours/days
- No 15-minute limit
- Up to 120GB RAM
- 4 vCPUs



MORE USE CASES

- Long-running processes
- Legacy applications
- Microservices
- Batch jobs and More...

Lambda vs Fargate - Decision Tree



Real Student Projects



ASSIGNMENT SUBMISSION PORTAL

- API endpoints (submit, list assignments)
- Email notifications
- File validation
- Cost: \$0.50/month for 500 students



CODE PLAGIARISM CHECKER

- Lambda: Receive submissions, trigger checks
- Fargate: Run complex similarity algorithms (30 min)
- Lambda: Send results email
- Cost: \$2/month for 200 submissions



LIVE STREAMING PLATFORM

- 24/7 streaming server
- Custom RTMP processing
- Could not use Lambda (continuous running)

Architecture Patterns to Try

- Event-Driven API (S3 upload → Lambda → DynamoDB)
- Scheduled Jobs (EventBridge (cron) → Lambda → Process data)
- Microservices (API Gateway → Lambda/Fargate → Services)
- Data Pipeline (S3 → Lambda (trigger) → Fargate (process) → S3)
- Webhook Handler (GitHub/Stripe → API Gateway → Lambda → Action)

Start small, scale infinitely!

The Serverless Mindset Shift

“I need a server to run my code”

- Provision capacity for peak load
- Manage infrastructure
- Pay for idle time
- Scale manually

The Serverless Mindset Shift

“I need to run code when X happens”

- Specify requirements
- Let AWS handle infrastructure
- Pay for execution only
- Scale automatically



Think in events, not servers!

When NOT to Use Serverless

“Be honest – serverless isn't always the answer”

- Predictable (and high) traffic
- When you need more control
- Long-running tasks
- When you aren't sure of your requirements

Key Takeaways

- Serverless = Ordering Pizza
- Lambda vs Fargate
- Focus on code, not servers
- Pay for value, not idle time



References

- <https://www.ben-morris.com/when-to-use-serverless-architecture-and-when-not-to/>
- <https://aws.amazon.com/blogs/compute/operating-lambda-anti-patterns-in-event-driven-architectures-part-3/>
- <https://aws.amazon.com/blogs/architecture/mistakes-to-avoid-when-implementing-serverless-architecture-with-lambda/>



Thank you!



Just search

The screenshot shows a Google search results page for the query "Avinash Dalvi AWS Community". The top result is a link to "Avinash Dalvi AWS Community Builder" with the URL <https://www.avinashdalvi.com>. The snippet below the link describes Avinash Shashikant Dalvi as a Full Stack Developer | PHP + Angular + Python + AWS | Speaker | Blogger | Leadership. Seasoned full-stack ...". The second result is a LinkedIn post from Avinash Dalvi on the Pulse platform, dated 9 months ago, with 10+ reactions. The snippet says: "Becoming an AWS Community Builder: It's About the Journey, Not the Badge Back in 2016, I started sharing my learnings about AWS—writing ...". The third result is another LinkedIn post from Avinash Dalvi, dated 4 months ago, with 130+ reactions. The snippet says: "Just received my AWS Community Builders swag pack — and I absolutely love it! Big thanks to Jason Dunn and the team for curating such ...". The fourth result is a tweet from Avinash Dalvi (@AvinashDalvi_) with over 1.1K followers. The snippet says: "Full Stack Developer | Speaker | Leadership | write blog at <https://t.co/mSjx5jlhD> Sr. Staff Engineer | ❤️ PHP & AWS.". The fifth result is a YouTube channel for Avinash Dalvi, with the URL <https://m.youtube.com/community>.